

TASK #1: UNDERSTAND THE PROBLEM STATEMENT AND BUSINESS CASE

- In this project, you have been hired as a data scientist at a bank and you have been provided with extensive data on the bank's customers for the past 6 months.
- Data includes transactions frequency, amount, tenure..etc.
- The bank marketing team would like to leverage AI/ML to launch a targeted marketing ad campaign that is tailored to specific group of customers.
- In order for this campaign to be successful, the bank has to divide its customers into at least 3 distinctive groups.
- This process is known as “marketing segmentation” and it crucial for maximizing marketing campaign conversion rate.



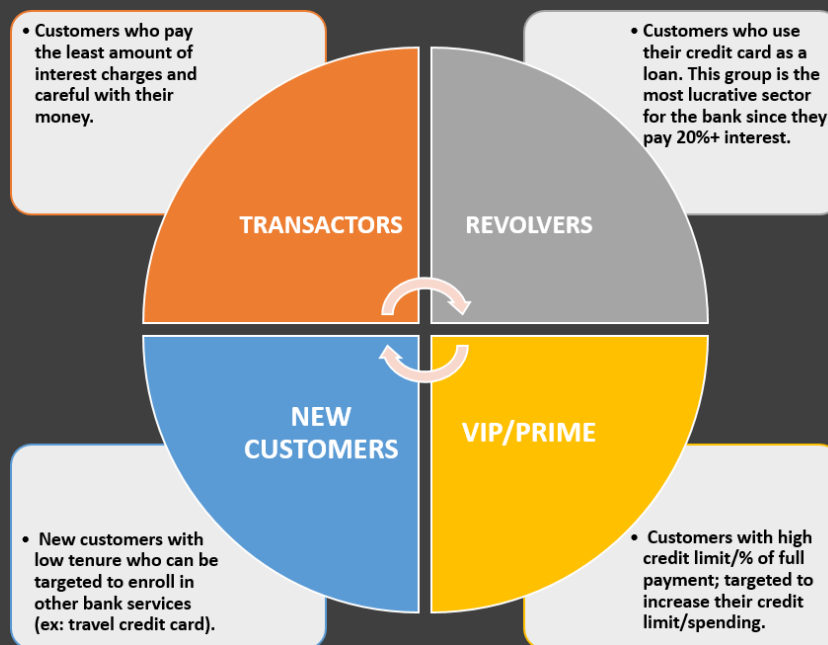
- Data Source: <https://www.kaggle.com/arjunbhasin2013/ccdata>
- Photo Credit: <https://www.needpix.com/photo/1011172/marketing-customer-polaroid-center-presentation-online-board-target-economy>

INSTRUCTOR

- Adjunct professor & online instructor
- Passionate about artificial intelligence, machine learning, and electric vehicles
- Taught 80,000+ students globally
- MBA (2018), Ph.D. (2014), M.A.Sc (2011)



Ryan Ahmed, Ph.D.



Data Source: <https://www.kaggle.com/arjunbhasin2013/ccdata> (<https://www.kaggle.com/arjunbhasin2013/ccdata>)

TASK #2: IMPORT LIBRARIES AND DATASETS

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler, normalize
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from jupyterthemes import jtplot
jtplot.style(theme='monokai', context='notebook', ticks=True, grid=False)
# setting the style of the notebook to be monokai theme
# this line of code is important to ensure that we are able to see the x and y axes clearly
# If you don't run this code line, you will notice that the xlabel and ylabel on any plot is black
on black and it will be hard to see them.
```

```
In [2]: # You have to include the full link to the csv file containing your dataset
creditcard_df = pd.read_csv('Marketing_data.csv')

# CUSTID: Identification of Credit Card holder
# BALANCE: Balance amount left in customer's account to make purchases
# BALANCE_FREQUENCY: How frequently the Balance is updated, score between 0 and 1 (1 = frequently
updated, 0 = not frequently updated)
# PURCHASES: Amount of purchases made from account
# ONEOFFPURCHASES: Maximum purchase amount done in one-go
# INSTALLMENTS_PURCHASES: Amount of purchase done in installment
# CASH_ADVANCE: Cash in advance given by the user
# PURCHASES_FREQUENCY: How frequently the Purchases are being made, score between 0 and 1 (1 = fre
quently purchased, 0 = not frequently purchased)
# ONEOFF_PURCHASES_FREQUENCY: How frequently Purchases are happening in one-go (1 = frequently pur
chased, 0 = not frequently purchased)
# PURCHASES_INSTALLMENTS_FREQUENCY: How frequently purchases in installments are being done (1 = f
requently done, 0 = not frequently done)
# CASH_ADVANCE_FREQUENCY: How frequently the cash in advance being paid
# CASH_ADVANCE_TRX: Number of Transactions made with "Cash in Advance"
# PURCHASES_TRX: Number of purchase transactions made
# CREDIT_LIMIT: Limit of Credit Card for user
# PAYMENTS: Amount of Payment done by user
# MINIMUM_PAYMENTS: Minimum amount of payments made by user
# PRC_FULL_PAYMENT: Percent of full payment paid by user
# TENURE: Tenure of credit card service for user
```

```
In [3]: creditcard_df
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CAS
0	C10001	40.900749	0.818182	95.40	0.00	95.40	0.000
1	C10002	3202.467416	0.909091	0.00	0.00	0.00	6442
2	C10003	2495.148862	1.000000	773.17	773.17	0.00	0.000
3	C10004	1666.670542	0.636364	1499.00	1499.00	0.00	205.7
4	C10005	817.714335	1.000000	16.00	16.00	0.00	0.000
...
8945	C19186	28.493517	1.000000	291.12	0.00	291.12	0.000
8946	C19187	19.183215	1.000000	300.00	0.00	300.00	0.000
8947	C19188	23.398673	0.833333	144.40	0.00	144.40	0.000
8948	C19189	13.457564	0.833333	0.00	0.00	0.00	36.55
8949	C19190	372.708075	0.666667	1093.25	1093.25	0.00	127.0

8950 rows x 18 columns

```
In [4]: creditcard_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8950 entries, 0 to 8949
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype  
---  -
0   CUST_ID                               8950 non-null   object  
1   BALANCE                               8950 non-null   float64 
2   BALANCE_FREQUENCY                     8950 non-null   float64 
3   PURCHASES                             8950 non-null   float64 
4   ONEOFF_PURCHASES                      8950 non-null   float64 
5   INSTALLMENTS_PURCHASES                8950 non-null   float64 
6   CASH_ADVANCE                          8950 non-null   float64 
7   PURCHASES_FREQUENCY                   8950 non-null   float64 
8   ONEOFF_PURCHASES_FREQUENCY            8950 non-null   float64 
9   PURCHASES_INSTALLMENTS_FREQUENCY      8950 non-null   float64 
10  CASH_ADVANCE_FREQUENCY                 8950 non-null   float64 
11  CASH_ADVANCE_TRX                       8950 non-null   int64   
12  PURCHASES_TRX                         8950 non-null   int64   
13  CREDIT_LIMIT                           8949 non-null   float64 
14  PAYMENTS                              8950 non-null   float64 
15  MINIMUM_PAYMENTS                      8637 non-null   float64 
16  PRC_FULL_PAYMENT                      8950 non-null   float64 
17  TENURE                                8950 non-null   int64   
dtypes: float64(14), int64(3), object(1)
memory usage: 1.2+ MB
```

```
In [ ]: # Let's apply info and get additional insights on our dataframe
        # 18 features with 8950 points
```

MINI CHALLENGE #1:

- What is the average, minimum and maximum "BALANCE" amount?

```
In [8]: print('Average, min, max =', creditcard_df['BALANCE'].mean(), creditcard_df['BALANCE'].min(), creditcard_df['BALANCE'].max())

Average, min, max = 1564.4748276781038 0.0 19043.13856
```

```
In [9]: creditcard_df.describe()

# Let's apply describe() and get more statistical insights on our dataframe
# Mean balance is $1564
# Balance frequency is frequently updated on average ~0.9
# Purchases average is $1000
# one off purchase average is ~$600
# Average purchases frequency is around 0.5
# average ONEOFF_PURCHASES_FREQUENCY, PURCHASES_INSTALLMENTS_FREQUENCY, and CASH_ADVANCE_FREQUENCY are generally low
# Average credit Limit ~ 4500
# Percent of full payment is 15%
# Average tenure is 11 years
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
count	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000
mean	1564.474828	0.877271	1003.204834	592.437371	411.067645	978.871112
std	2081.531879	0.236904	2136.634782	1659.887917	904.338115	2097.163877
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	128.281915	0.888889	39.635000	0.000000	0.000000	0.000000
50%	873.385231	1.000000	361.280000	38.000000	89.000000	0.000000
75%	2054.140036	1.000000	1110.130000	577.405000	468.637500	1113.821139
max	19043.138560	1.000000	49039.570000	40761.250000	22500.000000	47137.211760

MINI CHALLENGE #2:

- Obtain the features (row) of the customer who made the maximim "ONEOFF_PURCHASES"
- Obtain the features of the customer who made the maximum cash advance transaction? how many cash advance transactions did that customer make? how often did he/she pay their bill?

```
In [10]: creditcard_df[creditcard_df['ONEOFF_PURCHASES'] == 40761.25]
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH
550	C10574	11547.52001	1.0	49039.57	40761.25	8278.32	558.1€

```
In [11]: creditcard_df['CASH_ADVANCE'].max()
```

47137.211760000006

```
In [13]: creditcard_df[creditcard_df['CASH_ADVANCE'] == 47137.211760000006]
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH
2159	C12226	10905.05381	1.0	431.93	133.5	298.43	4713

TASK #3: VISUALIZE AND EXPLORE DATASET

```
In [14]: # Let's see if we have any missing data, luckily we don't have many!
sns.heatmap(creditcard_df.isnull(), yticklabels = False, cbar = False, cmap="Blues")
```

<matplotlib.axes._subplots.AxesSubplot at 0x24fb467a288>



```
In [15]: creditcard_df.isnull().sum()
```

```
CUST_ID          0
BALANCE          0
BALANCE_FREQUENCY 0
PURCHASES        0
ONEOFF_PURCHASES 0
INSTALLMENTS_PURCHASES 0
CASH_ADVANCE     0
PURCHASES_FREQUENCY 0
ONEOFF_PURCHASES_FREQUENCY 0
PURCHASES_INSTALLMENTS_FREQUENCY 0
CASH_ADVANCE_FREQUENCY 0
CASH_ADVANCE_TRX 0
PURCHASES_TRX    0
CREDIT_LIMIT     1
PAYMENTS         0
MINIMUM_PAYMENTS 313
PRC_FULL_PAYMENT 0
TENURE           0
dtype: int64
```

```
In [17]: # Fill up the missing elements with mean of the 'MINIMUM_PAYMENT'
creditcard_df.loc[(creditcard_df['MINIMUM_PAYMENTS'].isnull() == True), 'MINIMUM_PAYMENTS'] = creditcard_df['MINIMUM_PAYMENTS'].mean()
```

```
In [18]: creditcard_df.isnull().sum()
```

```
CUST_ID                0
BALANCE                0
BALANCE_FREQUENCY      0
PURCHASES              0
ONEOFF_PURCHASES       0
INSTALLMENTS_PURCHASES 0
CASH_ADVANCE           0
PURCHASES_FREQUENCY    0
ONEOFF_PURCHASES_FREQUENCY 0
PURCHASES_INSTALLMENTS_FREQUENCY 0
CASH_ADVANCE_FREQUENCY 0
CASH_ADVANCE_TRX       0
PURCHASES_TRX          0
CREDIT_LIMIT           1
PAYMENTS               0
MINIMUM_PAYMENTS       0
PRC_FULL_PAYMENT       0
TENURE                 0
dtype: int64
```

MINI CHALLENGE #3:

- Fill out missing elements in the "CREDIT_LIMIT" column
- Double check and make sure that no missing elements are present

```
In [19]: creditcard_df.loc[(creditcard_df['CREDIT_LIMIT'].isnull() == True), 'CREDIT_LIMIT'] = creditcard_d
f['CREDIT_LIMIT'].mean()
```

```
In [20]: creditcard_df.isnull().sum()
```

```
CUST_ID                0
BALANCE                0
BALANCE_FREQUENCY      0
PURCHASES              0
ONEOFF_PURCHASES       0
INSTALLMENTS_PURCHASES 0
CASH_ADVANCE           0
PURCHASES_FREQUENCY    0
ONEOFF_PURCHASES_FREQUENCY 0
PURCHASES_INSTALLMENTS_FREQUENCY 0
CASH_ADVANCE_FREQUENCY 0
CASH_ADVANCE_TRX       0
PURCHASES_TRX          0
CREDIT_LIMIT           0
PAYMENTS               0
MINIMUM_PAYMENTS       0
PRC_FULL_PAYMENT       0
TENURE                 0
dtype: int64
```



```
In [21]: sns.heatmap(creditcard_df.isnull(), yticklabels = False, cbar = False, cmap="Blues")
```

<matplotlib.axes._subplots.AxesSubplot at 0x24fb3047908>



```
In [22]: # Let's see if we have duplicated entries in the data
creditcard_df.duplicated().sum()
```

0

MINI CHALLENGE #4:

- Drop Customer ID column 'CUST_ID' and make sure that the column has been removed from the dataframe

```
In [23]: creditcard_df.drop('CUST_ID', axis = 1, inplace = True)
```

In [26]:

creditcard_df

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANC
0	40.900749	0.818182	95.40	0.00	95.40	0.000000
1	3202.467416	0.909091	0.00	0.00	0.00	6442.945483
2	2495.148862	1.000000	773.17	773.17	0.00	0.000000
3	1666.670542	0.636364	1499.00	1499.00	0.00	205.788017
4	817.714335	1.000000	16.00	16.00	0.00	0.000000
...
8945	28.493517	1.000000	291.12	0.00	291.12	0.000000
8946	19.183215	1.000000	300.00	0.00	300.00	0.000000
8947	23.398673	0.833333	144.40	0.00	144.40	0.000000
8948	13.457564	0.833333	0.00	0.00	0.00	36.558778
8949	372.708075	0.666667	1093.25	1093.25	0.00	127.040008

8950 rows x 17 columns

In [27]:

n = len(creditcard_df.columns)
n

17

In [28]:

creditcard_df.columns

Index(['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
 'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
 'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
 'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX',
 'CREDIT_LIMIT', 'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT',
 'TENURE'],
 dtype='object')

```
In [30]: # distplot combines the matplotlib.hist function with seaborn kdeplot()
# KDE Plot represents the Kernel Density Estimate
# KDE is used for visualizing the Probability Density of a continuous variable.
# KDE demonstrates the probability density at different values in a continuous variable.

# Mean of balance is $1500
# 'Balance_Frequency' for most customers is updated frequently ~1
# For 'PURCHASES_FREQUENCY', there are two distinct group of customers
# For 'ONEOFF_PURCHASES_FREQUENCY' and 'PURCHASES_INSTALLMENT_FREQUENCY' most users don't do one o
ff purchases or installment purchases frequently
# Very small number of customers pay their balance in full 'PRC_FULL_PAYMENT'~0
# Credit limit average is around $4500
# Most customers are ~11 years tenure

plt.figure(figsize = (10,50))
for i in range(len(creditcard_df.columns)):
    plt.subplot(17, 1, i+1)
    sns.distplot(creditcard_df[creditcard_df.columns[i]], kde_kws={"color": "b", "lw": 3, "label":
"KDE"}, hist_kws={"color": "g"})
    plt.title(creditcard_df.columns[i])

plt.tight_layout()
```

```

-----
ValueError                                Traceback (most recent call last)
~\anaconda31\lib\site-packages\statsmodels\nonparametric\kde.py in kdensityfft(X, kernel, bw, weights, gridsize, adjust, clip, cut, retgrid)
    450     try:
--> 451         bw = float(bw)
    452     except:

ValueError: could not convert string to float: 'scott'

During handling of the above exception, another exception occurred:

RuntimeError                                Traceback (most recent call last)
<ipython-input-30-314483609319> in <module>
    15 for i in range(len(creditcard_df.columns)):
    16     plt.subplot(17, 1, i+1)
--> 17     sns.distplot(creditcard_df[creditcard_df.columns[i]], kde_kws={"color": "b", "lw": 3, "label": "KDE"}, hist_kws={"color": "g"})
    18     plt.title(creditcard_df.columns[i])
    19

~\anaconda31\lib\site-packages\seaborn\distributions.py in distplot(a, bins, hist, kde, rug, fit, hist_kws, kde_kws, rug_kws, fit_kws, color, vertical, norm_hist, axlabel, label, ax)
    231     if kde:
    232         kde_color = kde_kws.pop("color", color)
--> 233         kdeplot(a, vertical=vertical, ax=ax, color=kde_color, **kde_kws)
    234         if kde_color != color:
    235             kde_kws["color"] = kde_color

~\anaconda31\lib\site-packages\seaborn\distributions.py in kdeplot(data, data2, shade, vertical, kernel, bw, gridsize, cut, clip, legend, cumulative, shade_lowest, cbar, cbar_ax, cbar_kws, ax, **kwargs)
    703     ax = _univariate_kdeplot(data, shade, vertical, kernel, bw,
    704                             gridsize, cut, clip, legend, ax,
--> 705                             cumulative=cumulative, **kwargs)
    706
    707     return ax

~\anaconda31\lib\site-packages\seaborn\distributions.py in _univariate_kdeplot(data, shade, vertical, kernel, bw, gridsize, cut, clip, legend, ax, cumulative, **kwargs)
    293     x, y = _statsmodels_univariate_kde(data, kernel, bw,
    294                                       gridsize, cut, clip,
--> 295                                       cumulative=cumulative)
    296     else:
    297         # Fall back to scipy if missing statsmodels

~\anaconda31\lib\site-packages\seaborn\distributions.py in _statsmodels_univariate_kde(data, kernel, bw, gridsize, cut, clip, cumulative)
    365     fft = kernel == "gau"
    366     kde = smnp.KDEUnivariate(data)
--> 367     kde.fit(kernel, bw, fft, gridsize=gridsize, cut=cut, clip=clip)
    368     if cumulative:
    369         grid, y = kde.support, kde.cdf

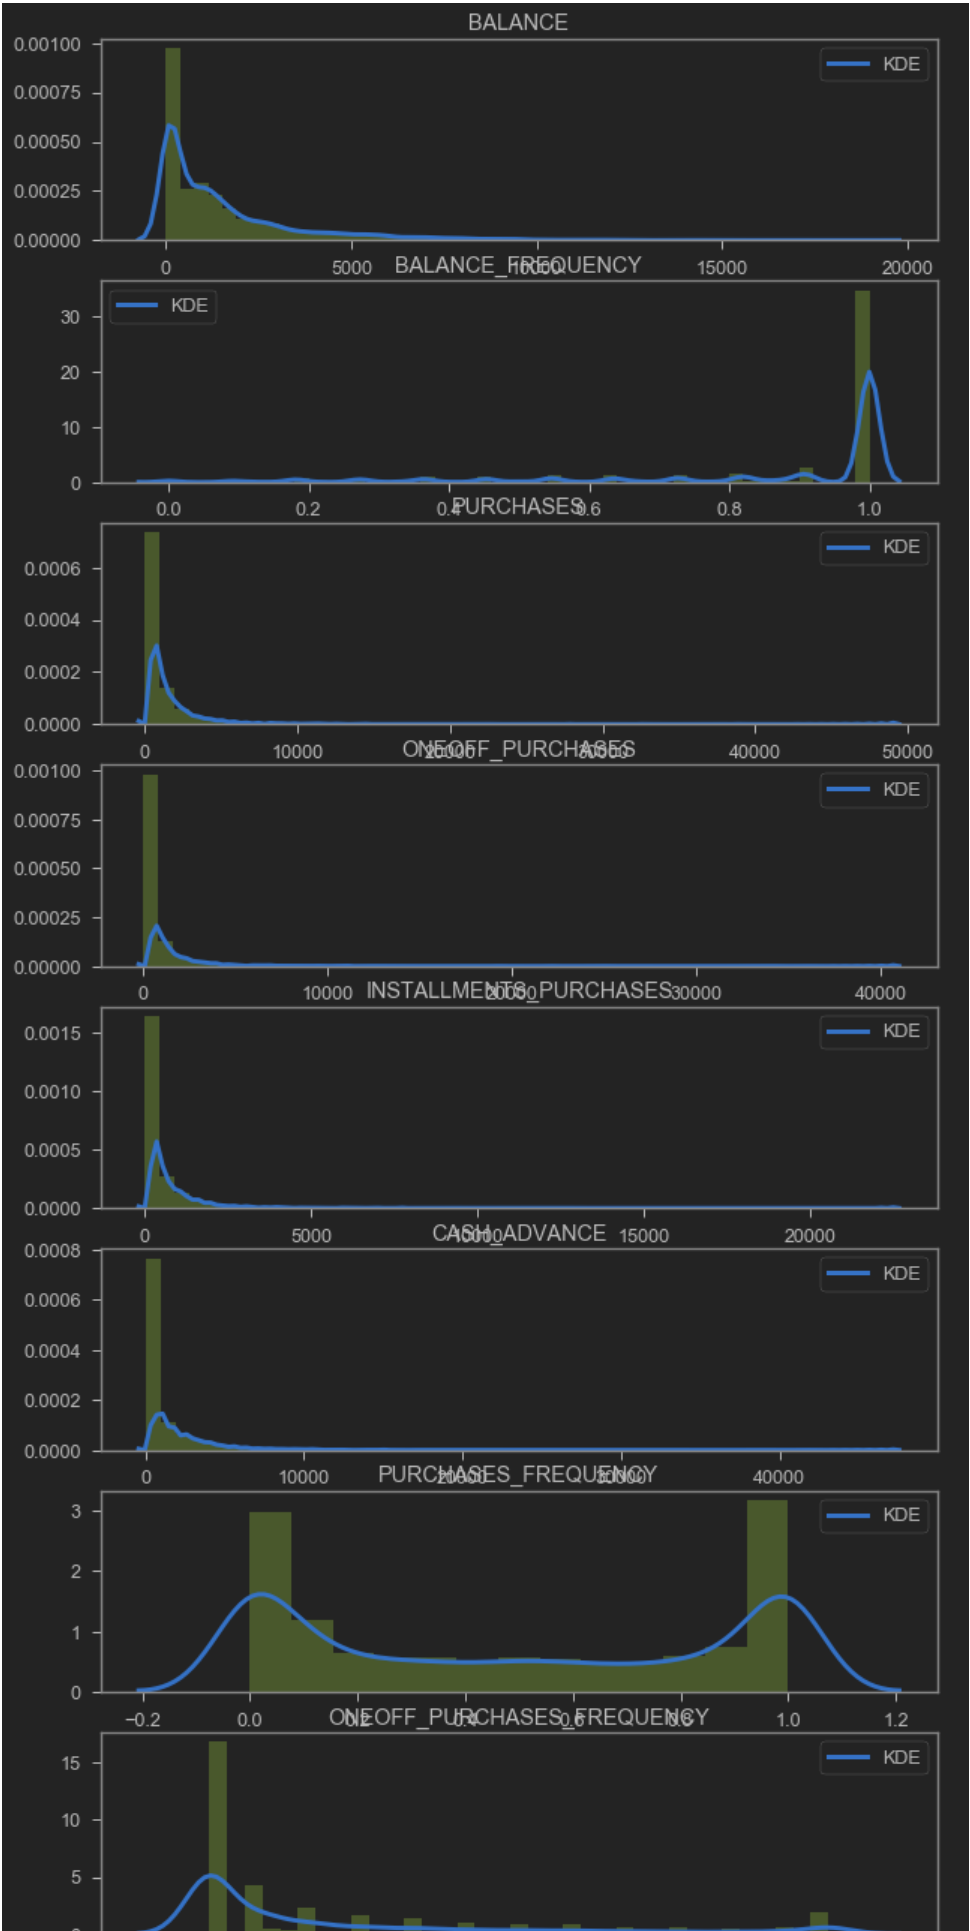
~\anaconda31\lib\site-packages\statsmodels\nonparametric\kde.py in fit(self, kernel, bw, fft, weights, gridsize, adjust, cut, clip)
    138     density, grid, bw = kdensityfft(endog, kernel=kernel, bw=bw,
    139                                     adjust=adjust, weights=weights, gridsize=gridsize,
--> 140                                     clip=clip, cut=cut)
    141     else:
    142         density, grid, bw = kdensity(endog, kernel=kernel, bw=bw,

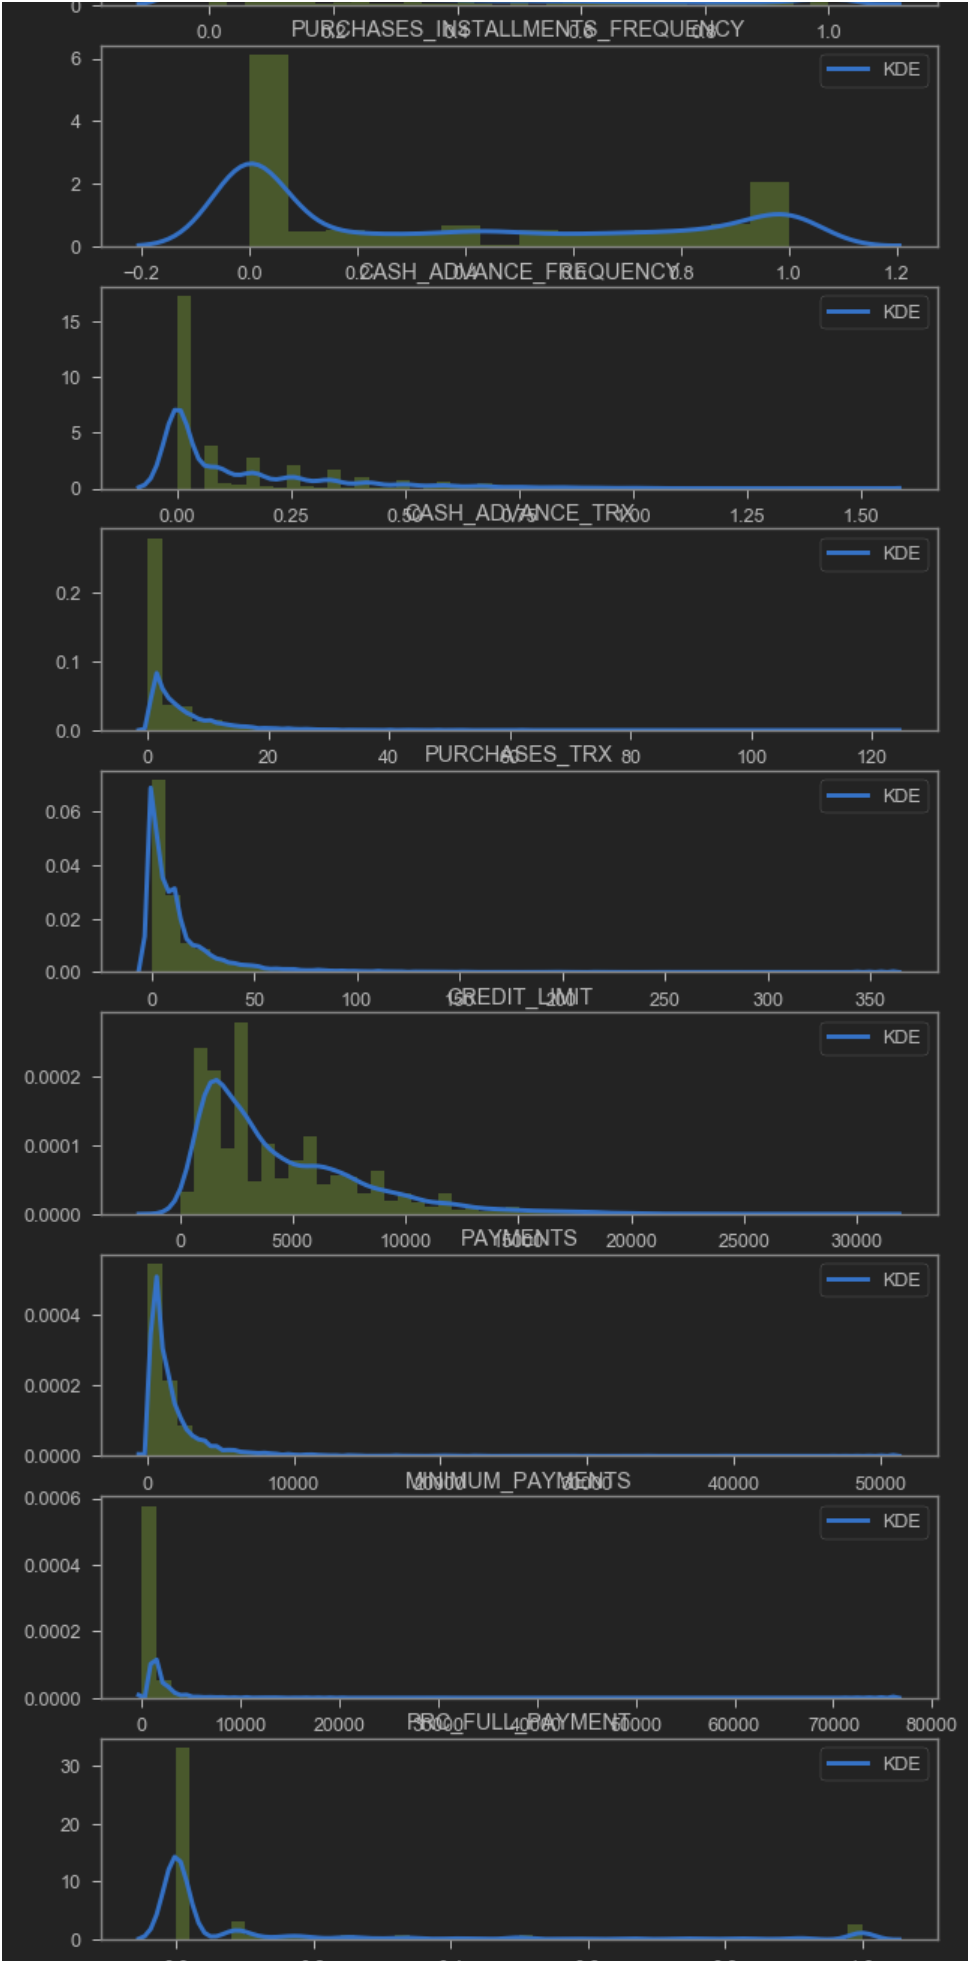
~\anaconda31\lib\site-packages\statsmodels\nonparametric\kde.py in kdensityfft(X, kernel, bw, weights, gridsize, adjust, clip, cut, retgrid)
    451     bw = float(bw)
    452     except:
--> 453     bw = bandwidths.select_bandwidth(X, bw, kern) # will cross-val fit this pattern?
    454     bw *= adjust
    455


~\anaconda31\lib\site-packages\statsmodels\nonparametric\bandwidths.py in select_bandwidth(x, bw, kernel)
    172     # eventually this can fall back on another selection criterion.
    173     err = "Selected KDE bandwidth is 0. Cannot estimate density."
--> 174     raise RuntimeError(err)
    175     else:
    176         return bandwidth

RuntimeError: Selected KDE bandwidth is 0. Cannot estimate density.

```





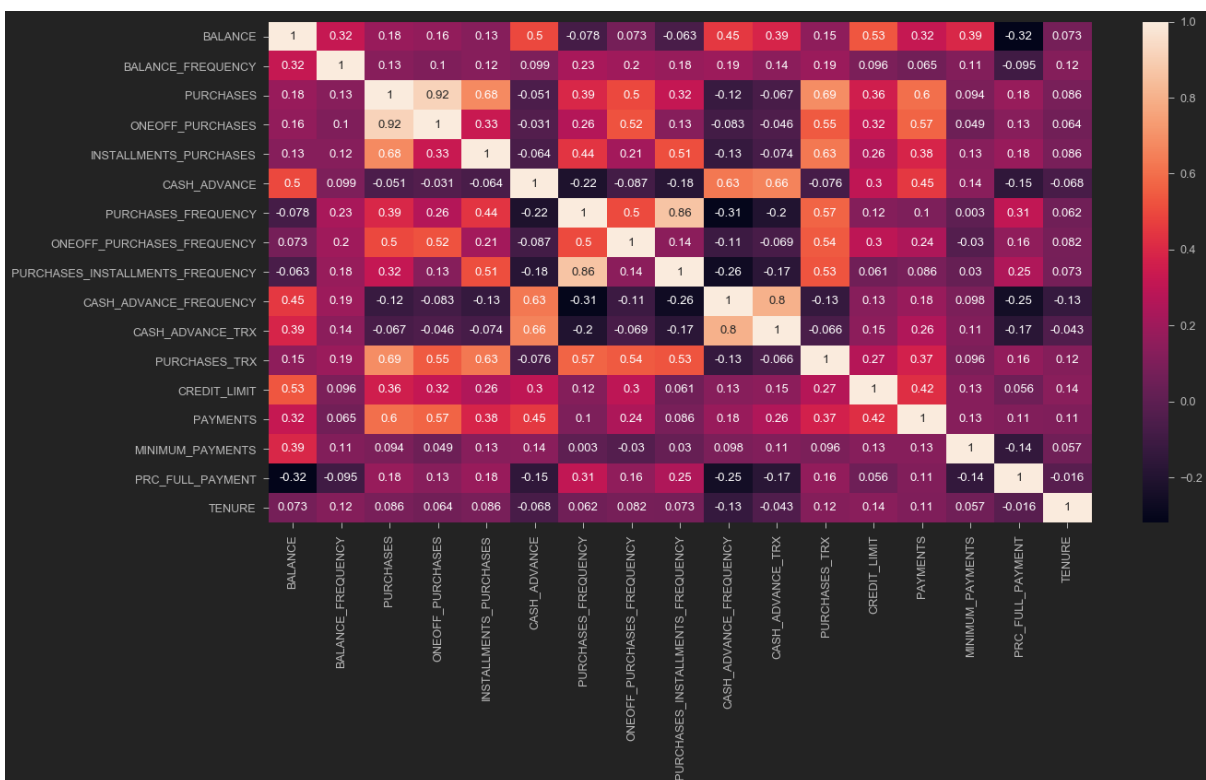


MINI CHALLENGE #5:

- Obtain the correlation matrix between features

```
In [32]: correlations = creditcard_df.corr()
f, ax = plt.subplots(figsize = (20, 10))
sns.heatmap(correlations, annot = True)
```


<matplotlib.axes._subplots.AxesSubplot at 0x24fb3c41688>



TASK #4: UNDERSTAND THE THEORY AND INTUITON BEHIND K-MEANS

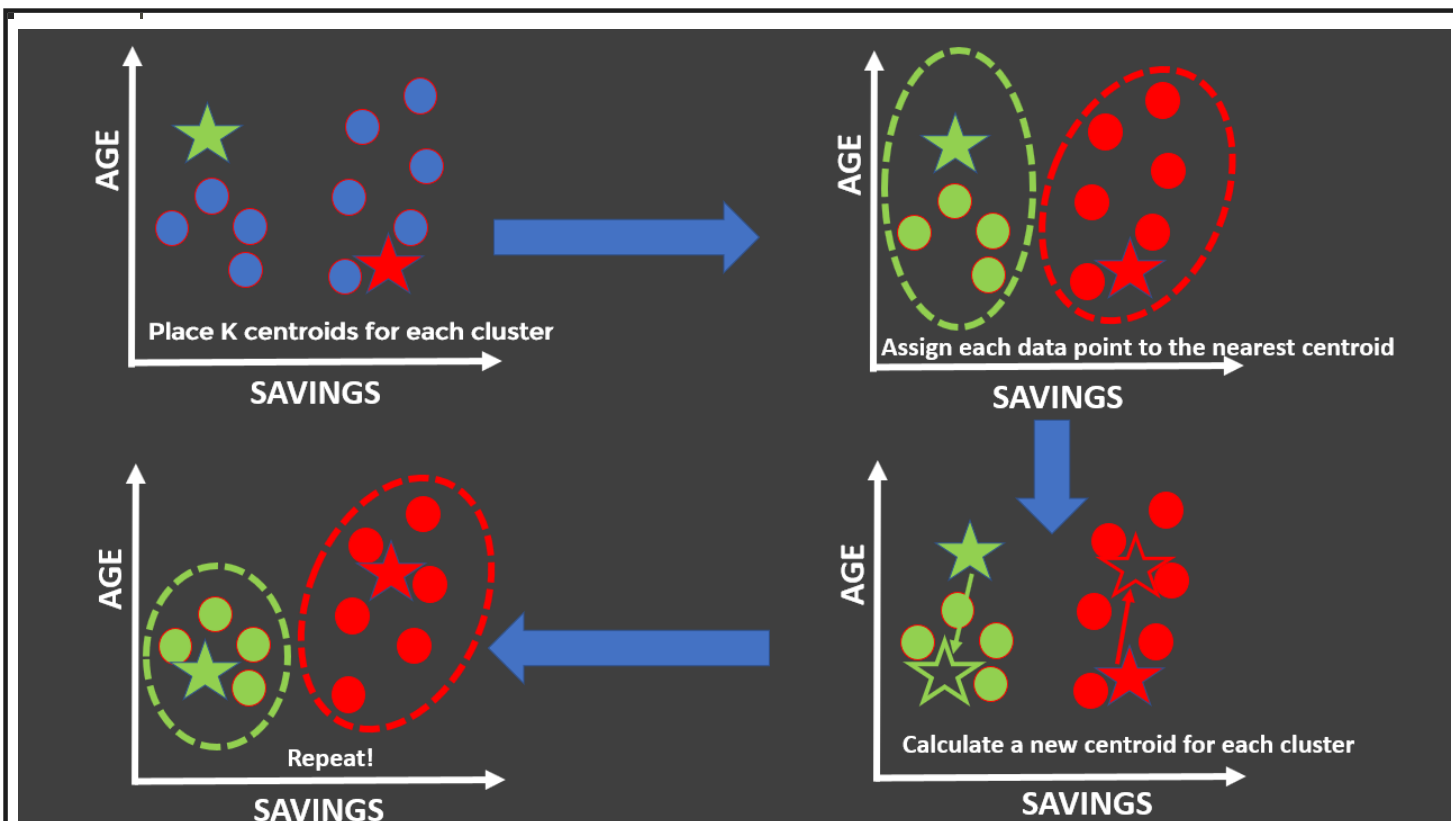
K-MEANS INTUITION

- K-means is an unsupervised learning algorithm (clustering).
- K-means works by grouping some data points together (clustering) in an unsupervised fashion.
- The algorithm groups observations with similar attribute values together by measuring the Euclidian distance between points.



K-MEANS ALGORITHM STEPS

1. Choose number of clusters “K”
2. Select random K points that are going to be the centroids for each cluster
3. Assign each data point to the nearest centroid, doing so will enable us to create “K” number of clusters
4. Calculate a new centroid for each cluster
5. Reassign each data point to the new closest centroid
6. Go to step 4 and repeat.



MINI CHALLENGE #6:

- Which of the following conditions could terminate the K-means clustering algorithm? (choose 2)
 - K-means terminates after a fixed number of iterations is reached
 - K-means terminates when the number of clusters does not increase between iterations
 - K-means terminates when the centroid locations do not change between iterations

```
In [ ]: // K-means terminates after a fixed number of iterations is reached
        // K-means terminates when the centroid locations do not change between iterations
```


TASK #5: LEARN HOW TO OBTAIN THE OPTIMAL NUMBER OF CLUSTERS (ELBOW METHOD)

HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

“ELBOW METHOD”

Within Cluster Sum of Squares (WCSS)

$$= \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$




HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

"ELBOW METHOD"

$$\text{Within Cluster Sum of Squares (WCSS)} = \sum_{P_i \text{ in Cluster } 1} \text{distance}(P_i, C_1)^2$$

NUMBER OF CLUSTERS (K) = 1




HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

“ELBOW METHOD”

$$\text{Within Cluster Sum of Squares (WCSS)} = \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2$$

NUMBER OF CLUSTERS (K) = 2




HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)?

“ELBOW METHOD”


Within Cluster Sum of Squares (WCSS)

$$= \sum_{P_i \text{ in Cluster 1}} \text{distance}(P_i, C_1)^2 + \sum_{P_i \text{ in Cluster 2}} \text{distance}(P_i, C_2)^2 + \sum_{P_i \text{ in Cluster 3}} \text{distance}(P_i, C_3)^2$$

NUMBER OF CLUSTERS (K) = 3



HOW TO SELECT THE OPTIMAL NUMBER OF CLUSTERS (K)? “ELBOW METHOD”



Source: https://commons.wikimedia.org/wiki/File:Tennis_Elbow_Illustration.jpg

TASK #6: FIND THE OPTIMAL NUMBER OF CLUSTERS USING ELBOW METHOD

- The elbow method is a heuristic method of interpretation and validation of consistency within cluster analysis designed to help find the appropriate number of clusters in a dataset.
- If the line chart looks like an arm, then the "elbow" on the arm is the value of k that is the best.
- Source:
 - [https://en.wikipedia.org/wiki/Elbow_method_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering)) ([https://en.wikipedia.org/wiki/Elbow_method_\(clustering\)](https://en.wikipedia.org/wiki/Elbow_method_(clustering)))
 - <https://www.geeksforgeeks.org/elbow-method-for-optimal-value-of-k-in-kmeans/> (<https://www.geeksforgeeks.org/elbow-method-for-optimal-value-of-k-in-kmeans/>)

```
In [33]: # Let's scale the data first
scaler = StandardScaler()
creditcard_df_scaled = scaler.fit_transform(creditcard_df)
```

```
In [34]: creditcard_df_scaled.shape

(8950, 17)
```

In [38]: creditcard_df_scaled

```
array([[ -0.73198937, -0.24943448, -0.42489974, ..., -0.31096755,
        -0.52555097,  0.36067954],
       [ 0.78696085,  0.13432467, -0.46955188, ...,  0.08931021,
        0.2342269 ,  0.36067954],
       [ 0.44713513,  0.51808382, -0.10766823, ..., -0.10166318,
        -0.52555097,  0.36067954],
       ...,
       [-0.7403981 , -0.18547673, -0.40196519, ..., -0.33546549,
        0.32919999, -4.12276757],
       [-0.74517423, -0.18547673, -0.46955188, ..., -0.34690648,
        0.32919999, -4.12276757],
       [-0.57257511, -0.88903307,  0.04214581, ..., -0.33294642,
        -0.52555097, -4.12276757]])
```

In [39]:

```
# Index(['BALANCE', 'BALANCE_FREQUENCY', 'PURCHASES', 'ONEOFF_PURCHASES',
#       'INSTALLMENTS_PURCHASES', 'CASH_ADVANCE', 'PURCHASES_FREQUENCY',
#       'ONEOFF_PURCHASES_FREQUENCY', 'PURCHASES_INSTALLMENTS_FREQUENCY',
#       'CASH_ADVANCE_FREQUENCY', 'CASH_ADVANCE_TRX', 'PURCHASES_TRX',
#       'CREDIT_LIMIT', 'PAYMENTS', 'MINIMUM_PAYMENTS', 'PRC_FULL_PAYMENT',
#       'TENURE'], dtype='object')
```


```
scores_1 = []
range_values = range(1, 20)
```

```
for i in range_values:
    kmeans = KMeans(n_clusters = i)
    kmeans.fit(creditcard_df_scaled)
    scores_1.append(kmeans.inertia_)
```

```
plt.plot(scores_1, 'bx-')
```

```
# From this we can observe that, 4th cluster seems to be forming the elbow of the curve.
# However, the values does not reduce linearly until 8th cluster.
# Let's choose the number of clusters to be 7 or 8.
```

```
[<matplotlib.lines.Line2D at 0x24fb3b82308>]
```




MINI CHALLENGE #7:

- Let's assume that our data only consists of the first 7 columns of "creditcard_df_scaled", what is the optimal number of clusters would be in this case? modify the code and rerun the cells.

```
In [43]: creditcard_df_scaled[:, :7].shape  
  
(8950, 7)
```

```
In [42]: scores_1 = []  
range_values = range(1, 20)  
  
for i in range_values:  
    kmeans = KMeans(n_clusters = i)  
    kmeans.fit(creditcard_df_scaled[:, :7])  
    scores_1.append(kmeans.inertia_)  
  
plt.plot(scores_1, 'bx-')
```

[<matplotlib.lines.Line2D at 0x24fb5e25888>]



TASK #7: APPLY K-MEANS METHOD

```
In [44]: kmeans = KMeans(7)  
kmeans.fit(creditcard_df_scaled)  
labels = kmeans.labels_ # Labels (cluster) associated to each data point
```

```
In [46]: kmeans.cluster_centers_.shape  
  
(7, 17)
```

```
In [47]: cluster_centers = pd.DataFrame(data = kmeans.cluster_centers_, columns = [creditcard_df.columns])
cluster_centers
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	0.009508	0.401739	-0.343430	-0.223474	-0.401364	-0.101925
1	-0.368878	0.333102	-0.041519	-0.230904	0.325868	-0.367172
2	1.488505	0.403475	7.413638	6.553369	5.486972	0.028557
3	-0.335429	-0.342203	-0.283935	-0.208973	-0.287081	0.064839
4	0.143949	0.430997	0.976283	0.923928	0.610965	-0.306762
5	1.675303	0.393689	-0.196573	-0.147680	-0.193575	1.996838
6	-0.701924	-2.134261	-0.306883	-0.230464	-0.302103	-0.322950

```
In [48]: # In order to understand what these numbers mean, Let's perform inverse transformation
cluster_centers = scaler.inverse_transform(cluster_centers)
cluster_centers = pd.DataFrame(data = cluster_centers, columns = [creditcard_df.columns])
cluster_centers
```

First Customers cluster (Transactors): Those are customers who pay least amount of intrerest charges and careful with their money, Cluster with lowest balance (\$104) and cash advance (\$303), Percentage of full payment = 23%

Second customers cluster (revolvers) who use credit card as a loan (most lucrative sector): high est balance (\$5000) and cash advance (~\$5000), Low purchase frequency, high cash advance frequency (0.5), high cash advance transactions (16) and low percentage of full payment (3%)

Third customer cluster (VIP/Prime): high credit limit \$16K and highest percentage of full payment, target for increase credit limit and increase spending habits

Fourth customer cluster (Low tenure): these are customers with low tenure (7 years), Low balance

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	1584.264568	0.972439	269.461234	221.516190	48.119482	765.130009
1	796.685924	0.956179	914.498711	209.184157	705.745881	208.893586
2	4662.671853	0.972850	16842.556892	11469.688108	5372.868784	1038.757441
3	866.307935	0.796206	396.572925	245.585564	151.464308	1114.841673
4	1864.092559	0.979370	3089.048627	2125.968921	963.555897	335.575951
5	5051.477573	0.970532	583.224191	347.318663	236.019764	5166.333011
6	103.479822	0.371684	347.544601	209.914180	137.880042	301.630330

```
In [49]: labels.shape # Labels associated to each data point

(8950,)
```

```
In [50]: labels.max()

6
```

```
In [51]: labels.min()

0
```

```
In [52]: y_kmeans = kmeans.fit_predict(creditcard_df_scaled)
y_kmeans

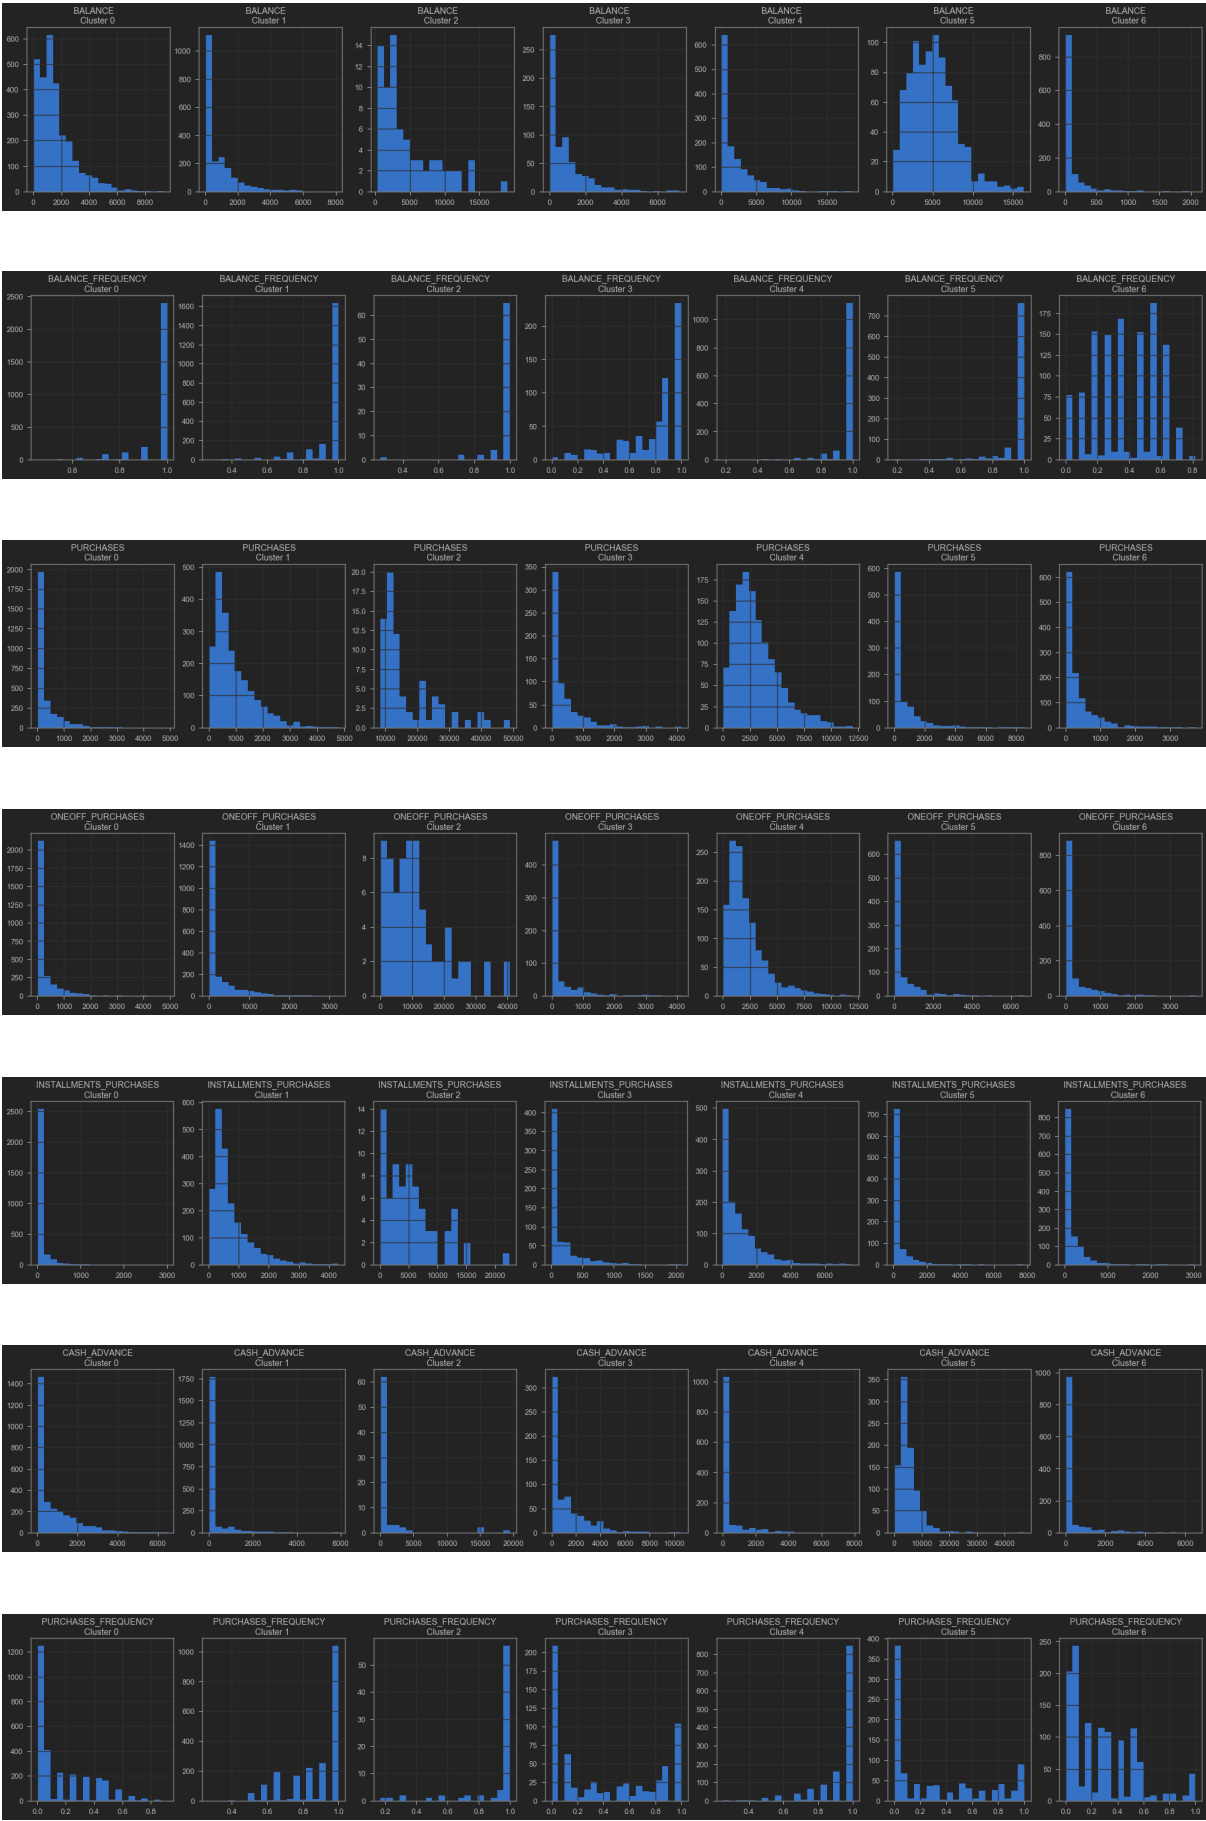
array([3, 6, 1, ..., 0, 3, 4])
```

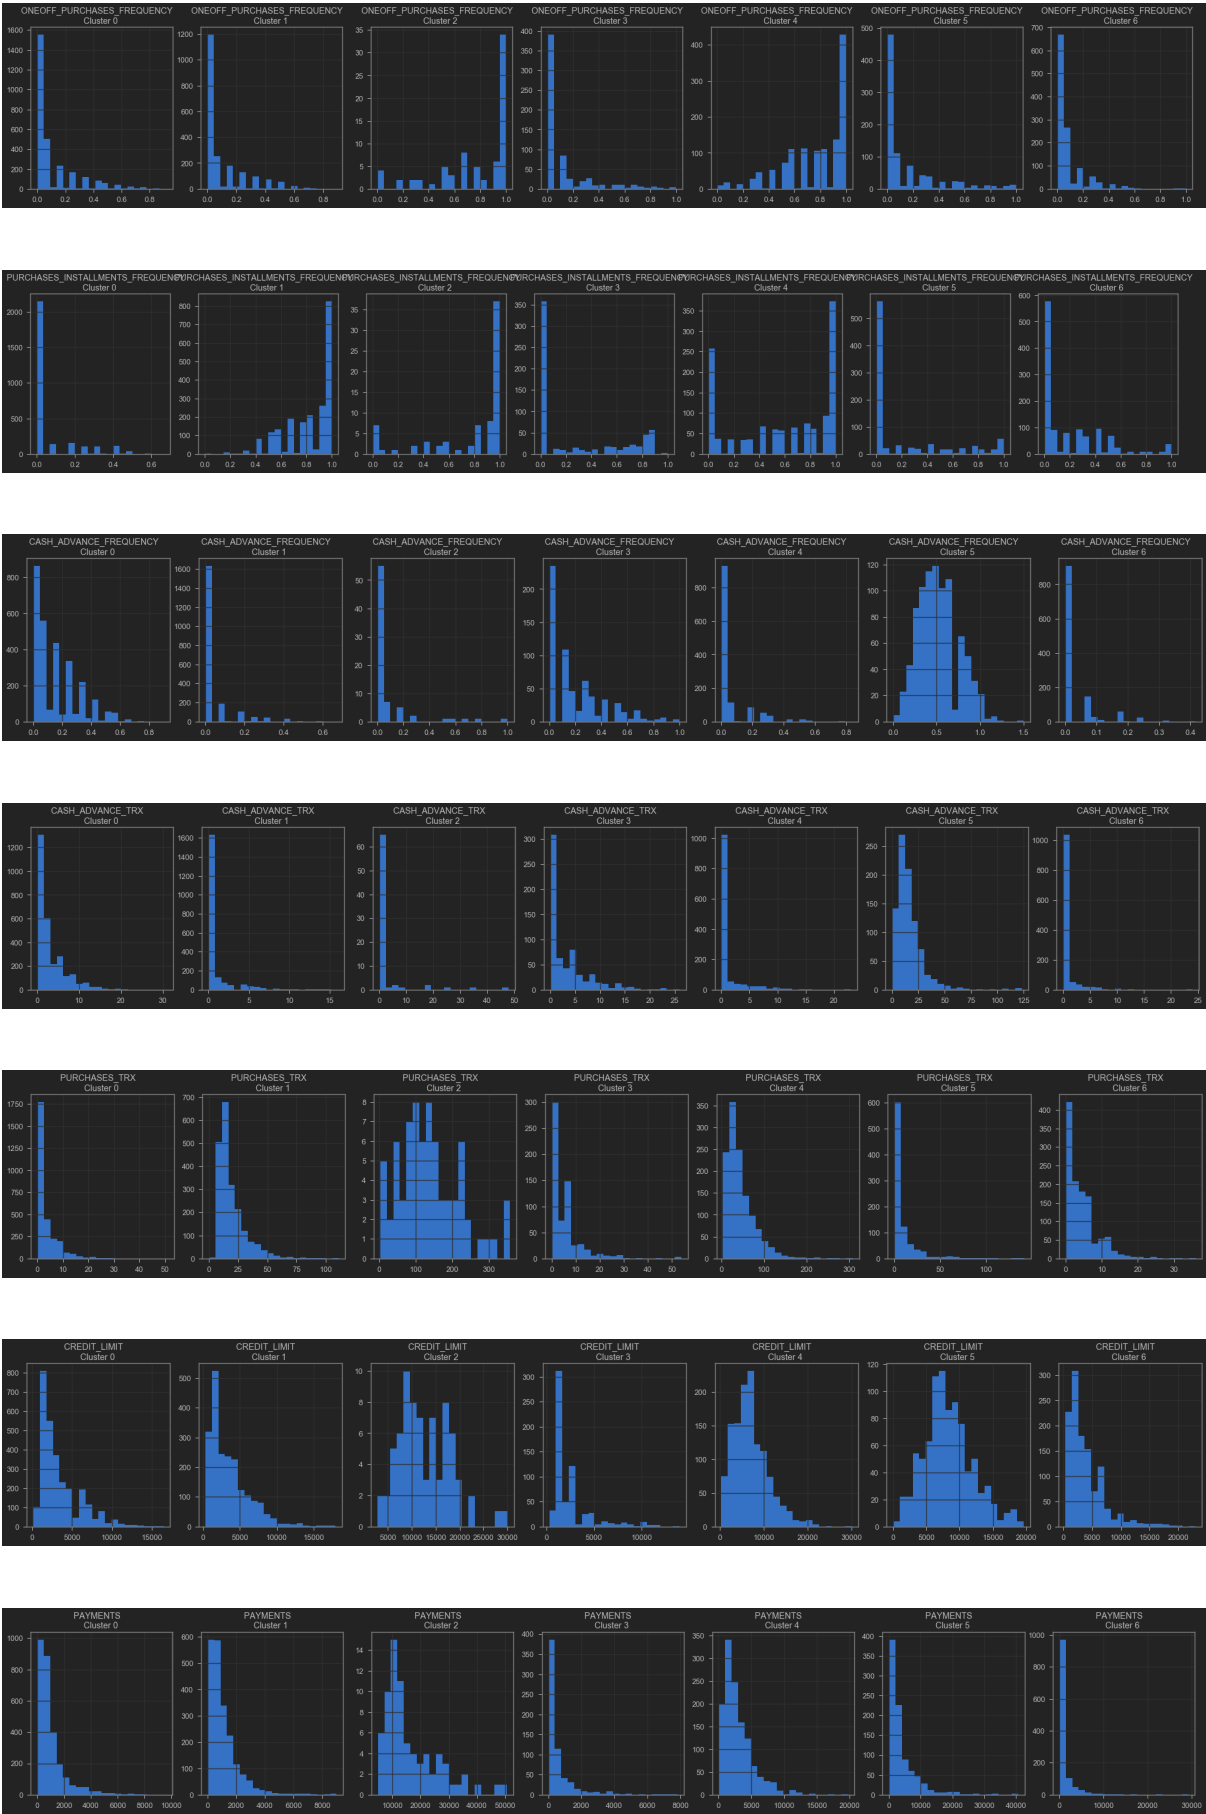
```
In [53]: # concatenate the clusters labels to our original dataframe
creditcard_df_cluster = pd.concat([creditcard_df, pd.DataFrame({'cluster':labels})], axis = 1)
creditcard_df_cluster.head()
```

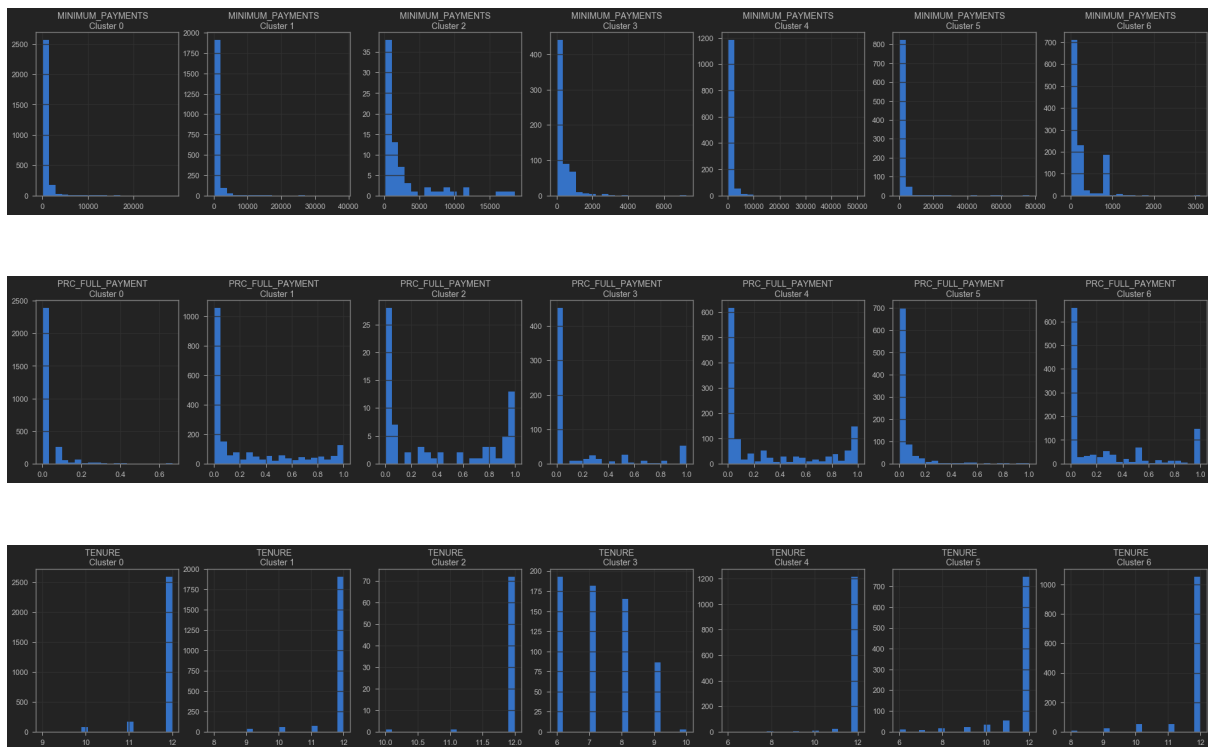
	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE
0	40.900749	0.818182	95.40	0.00	95.4	0.000000
1	3202.467416	0.909091	0.00	0.00	0.0	6442.945483
2	2495.148862	1.000000	773.17	773.17	0.0	0.000000
3	1666.670542	0.636364	1499.00	1499.00	0.0	205.788017
4	817.714335	1.000000	16.00	16.00	0.0	0.000000

```
In [54]: # Plot the histogram of various clusters
for i in creditcard_df.columns:
    plt.figure(figsize = (35, 5))
    for j in range(7):
        plt.subplot(1,7,j+1)
        cluster = creditcard_df_cluster[creditcard_df_cluster['cluster'] == j]
        cluster[i].hist(bins = 20)
        plt.title('{} \nCluster {}'.format(i,j))

plt.show()
```







MINI CHALLENGE #8:

- Repeat the same procedure with 8 clusters instead of 7

TASK 8: APPLY PRINCIPAL COMPONENT ANALYSIS AND VISUALIZE THE RESULTS

PRINCIPAL COMPONENT ANALYSIS (PCA)

- PCA is an unsupervised machine learning algorithm.
- PCA performs dimensionality reductions while attempting at keeping the original information unchanged.
- PCA works by trying to find a new set of features called components.
- Components are composites of the uncorrelated given input features.

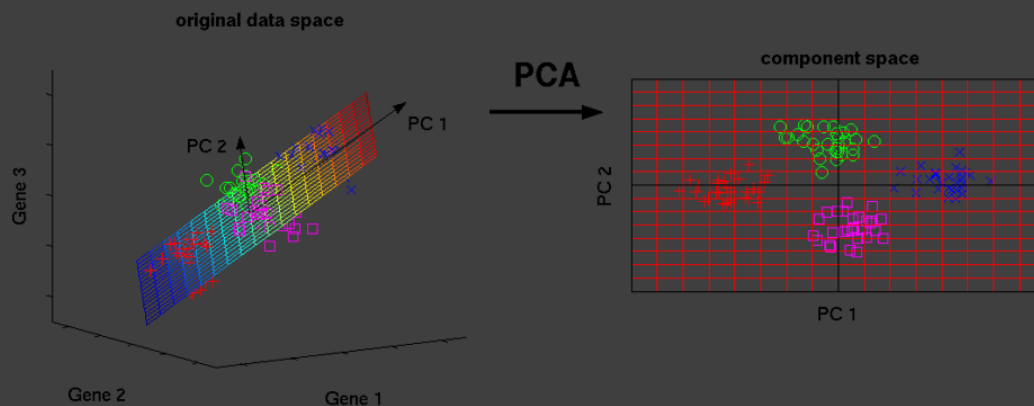


Photo Credit: <http://phdthesis-bioinformatics-maxplanckinstitute-molecularplantphys.matthias-scholz.de/>

```
In [56]: # Obtain the principal components
pca = PCA(n_components = 2)
principal_comp = pca.fit_transform(creditcard_df_scaled)
principal_comp

array([[ -1.68222052, -1.0764523 ],
       [ -1.13829595,  2.50646726],
       [  0.96968604, -0.38350322],
       ...,
       [ -0.92620362, -1.81078458],
       [ -2.33655245, -0.65797202],
       [ -0.55642189, -0.40046605]])
```




```
In [57]: # Create a dataframe with the two components
pca_df = pd.DataFrame(data = principal_comp, columns = ['pca1', 'pca2'])
pca_df.head()
```

	pca1	pca2
0	-1.682221	-1.076452
1	-1.138296	2.506467
2	0.969686	-0.383503
3	-0.873628	0.043164
4	-1.599434	-0.688581

```
In [58]: # Concatenate the clusters labels to the dataframe
pca_df = pd.concat([pca_df, pd.DataFrame({'cluster': labels})], axis = 1)
pca_df.head()
```

	pca1	pca2	cluster
0	-1.682221	-1.076452	0
1	-1.138296	2.506467	5
2	0.969686	-0.383503	4
3	-0.873628	0.043164	0
4	-1.599434	-0.688581	0

```
In [61]: plt.figure(figsize=(10,10))
ax = sns.scatterplot(x="pca1", y="pca2", hue = "cluster", data = pca_df, palette = ['red', 'pink', 'blue', 'green', 'purple', 'gray', 'yellow'])
plt.show()
```



MINI CHALLENGE #9:

- Repeat task #7 and #8 with number of clusters = 7 and 4

EXCELLENT JOB! YOU SHOULD BE PROUD OF YOUR NEWLY ACQUIRED SKILLS

MINI CHALLENGE SOLUTIONS

MINI CHALLENGE #1

```
In [ ]: # Average, minimum and maximum balance amounts
print('The average, minimum and maximum balance amount are:', creditcard_df['BALANCE'].mean(), creditcard_df['BALANCE'].min(), creditcard_df['BALANCE'].max())
```

MINI CHALLENGE #2

```
In [ ]: # Let's see who made one off purchase of $40761!
creditcard_df[creditcard_df['ONEOFF_PURCHASES'] == 40761.25]
```

```
In [ ]: creditcard_df['CASH_ADVANCE'].max()
```

```
In [ ]: # Let's see who made cash advance of $47137!
# This customer made 123 cash advance transactions!!
# Never paid credit card in full

creditcard_df[creditcard_df['CASH_ADVANCE'] == 47137.211760000006]
```

MINI CHALLENGE #3

```
In [ ]: # Fill up the missing elements with mean of the 'CREDIT_LIMIT'
creditcard_df.loc[(creditcard_df['CREDIT_LIMIT'].isnull() == True), 'CREDIT_LIMIT'] = creditcard_df['CREDIT_LIMIT'].mean()
sns.heatmap(creditcard_df.isnull(), yticklabels = False, cbar = False, cmap="Blues")
```

MINI CHALLENGE #4

```
In [ ]: # Let's drop Customer ID since it has no meaning here
creditcard_df.drop("CUST_ID", axis = 1, inplace= True)
creditcard_df.head()
```

MINI CHALLENGE #5

```
In [ ]: correlations = creditcard_df.corr()
f, ax = plt.subplots(figsize = (20, 20))
sns.heatmap(correlations, annot = True)

# 'PURCHASES' have high correlation between one-off purchases, 'installment purchases, purchase transactions, credit limit and payments.
# Strong Positive Correlation between 'PURCHASES_FREQUENCY' and 'PURCHASES_INSTALLMENT_FREQUENCY'
```

MINI CHALLENGE #6:

- Which of the following conditions could terminate the K-means clustering algorithm? (choose 2)
 - K-means terminates after a fixed number of iterations is reached (True)
 - K-means terminates when the number of clusters does not increase between iterations (False)
 - K-means terminates when the centroid locations do not change between iterations (True)

MINI CHALLENGE #7:

```
In [ ]: # code modification
        kmeans.fit(creditcard_df_scaled[:, :7])
        # optimal number of clusters would be = 5
```

MINI CHALLENGE #8 & #9:

- simply change the values requested in the question and rerun the cells